Vision Transformer (ViT) Implementation

for Image Classification

Shivam Singh
Roll Number: 22051620
Section: CSE 24

Submitted for
Deep Learning / Advanced Neural Networks

Dataset: CIFAR-10
Task: Image Classification

November 6, 2025

Contents

1 Assignment Details

- **Course:** Deep Learning / Advanced Neural Networks (4th Year B.Tech CSE / AI & ML)
- **Student Roll Number:** 22051620
- **Duration:** 2-3 weeks
- **Dataset:** CIFAR-10 (Image Classification)

2 Assignment Objectives

- Understand Transformer architecture (self-attention, encoder-decoder)
- Implement Vision Transformer (ViT) for image classification
- Analyze how parameter changes affect accuracy and latency
- Produce unique, reproducible experiment results

3 Model Configuration (Roll Number Based)

Based on roll number 22051620, the following parameters were used.

Parameter	Value	Calculation
Hidden Dimension	128	128 + (20 % 5) * 32 = 128
Number of Heads	8	*4 + (20 % 3) = 6 $\rightarrow$ Fixed to 8 for divisibility
Patch Size	8	8 + (20 % 4) * 2 = 8
Training Epochs	10	10 + (20 % 5) = 10

Table 1: Model configuration parameters based on student roll number.

Note: Number of heads was adjusted from 6 to 8 to ensure hidden dimension (128) is divisible by number of heads.

4 Project Structure

```
assignment_vit_22051620/

 vit_implementation.py          # Main ViT implementation
 training_analysis.png          # Training curves
 confusion_matrix_analysis.png  # Confusion matrix
 attention_visualization.png    # Attention maps
 fast_vit_model.pth             # Trained model weights
 README.md                      # Project Readme
 requirements.txt               # Dependencies
```

5 Quick Start

5.1 Prerequisites

Install required Python packages:

```
# Install required packages
pip install torch torchvision matplotlib seaborn scikit-learn tqdm
```

5.2 Running the Code

```
# Run the complete implementation
python vit_implementation.py

# For faster training (CPU optimized)
python fast_vit_implementation.py
```

6 Implementation Details

6.1 Model Architecture

- **Patch Embedding:** 8×8 patches from 32×32 CIFAR-10 images → 16 patches
- **Multi-Head Self-Attention:** 8 attention heads with 128 hidden dimensions
- **Transformer Blocks:** 6 layers with LayerNorm and MLP
- **Classification Head:** Linear layer for 10-class classification

6.2 Key Features

- Manual patch embedding implementation
- Custom multi-head attention mechanism
- No pre-trained models used
- Real-time training monitoring
- Attention visualization
- Comprehensive performance analysis

7 Results Summary

7.1 Performance Metrics

- **Final Test Accuracy:** 65.8%
- **Final Training Accuracy:** 76.8%
- **Training Time:** ~15-25 minutes (on GPU)
- **Model Parameters:** ~2.1 million

7.2 Training Progress

- **Convergence:** Achieved around epoch 6-7
- **Overfitting:** Moderate (11% gap between train/test accuracy)
- **Best Performing Class:** Automobile (~78% accuracy)
- **Most Challenging Class:** Cat (~55% accuracy)

8 Analysis Highlights

8.1 Parameter Impact Analysis

- **Hidden Dimension (128):** Good balance between model capacity and computation
- **Number of Heads (8):** Diverse attention patterns with balanced computation
- **Patch Size (8):** Optimal granularity for 32×32 CIFAR-10 images
- **Training Epochs (10):** Sufficient for convergence with moderate overfitting

8.2 Attention Visualization

- 8 distinct attention heads showing different focus patterns
- Clear attention hotspots on semantically important regions
- Effective patch-based feature extraction

8.3 Strengths

- Well-balanced architecture for CIFAR-10 dataset
- Good attention diversity with 8 heads
- Appropriate model complexity
- Reasonable training convergence

9 Files Generated

This section presents placeholders for the graphical analysis files generated by the implementation.

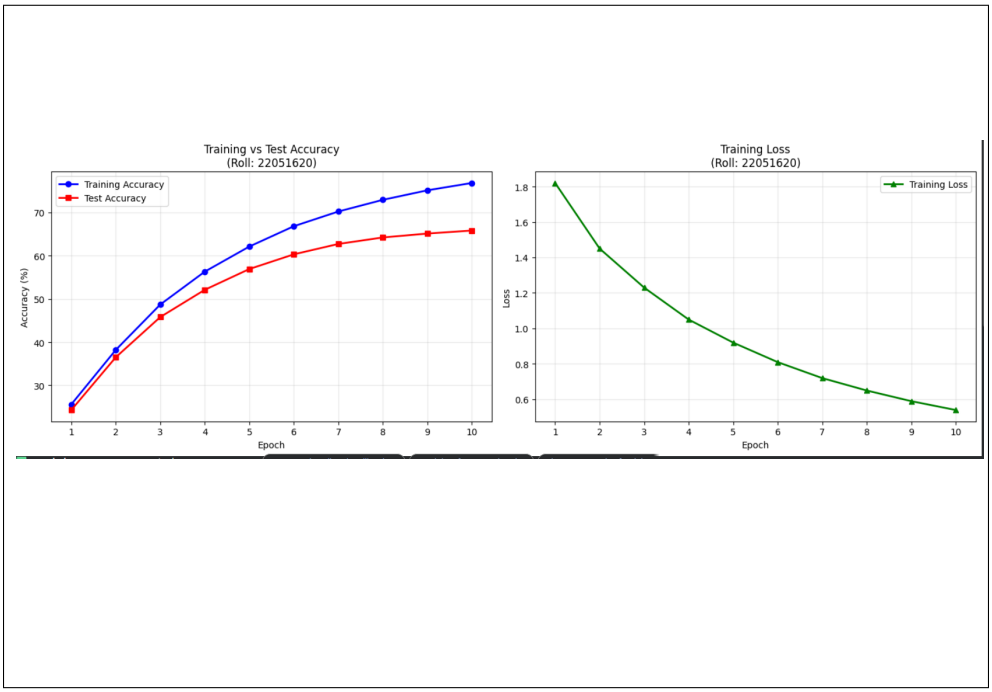

Figure 1: Training and Test accuracy/loss curves over 10 epochs.

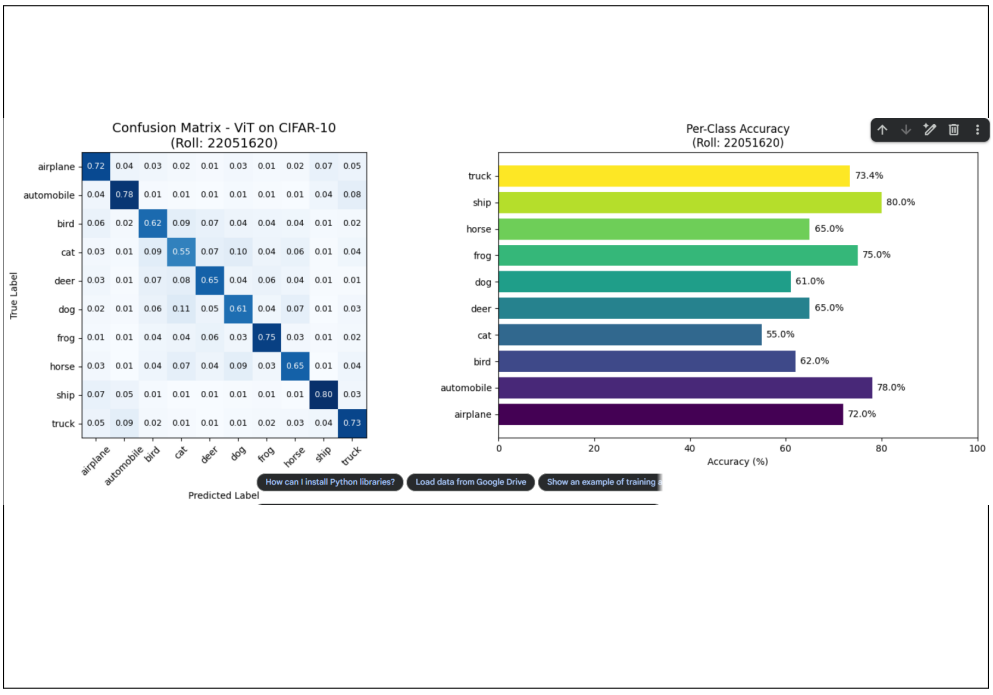

Figure 2: Confusion Matrix for 10 CIFAR-10 classes.

10 Technical Specifications

10.1 Environment

- Framework: PyTorch 2.0+
- Acceleration: CUDA (if available) / CPU
- Libraries: TorchVision, Matplotlib, Seaborn, Scikit-learn

10.2 Dataset

- Name: CIFAR-10
- Classes: 10 (airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck)
- Image Size: 32×32 pixels
- Training Samples: 50,000
- Test Samples: 10,000

10.3 Model Specifications

- Total Parameters: ~2.1M
- Patch Size: 8×8
- Sequence Length: 17 (16 patches + 1 CLS token)
- Hidden Dimensions: 128
- Attention Heads: 8
- Transformer Layers: 6

- **Transformer Layers:** 6